

# Engineering of ML Systems

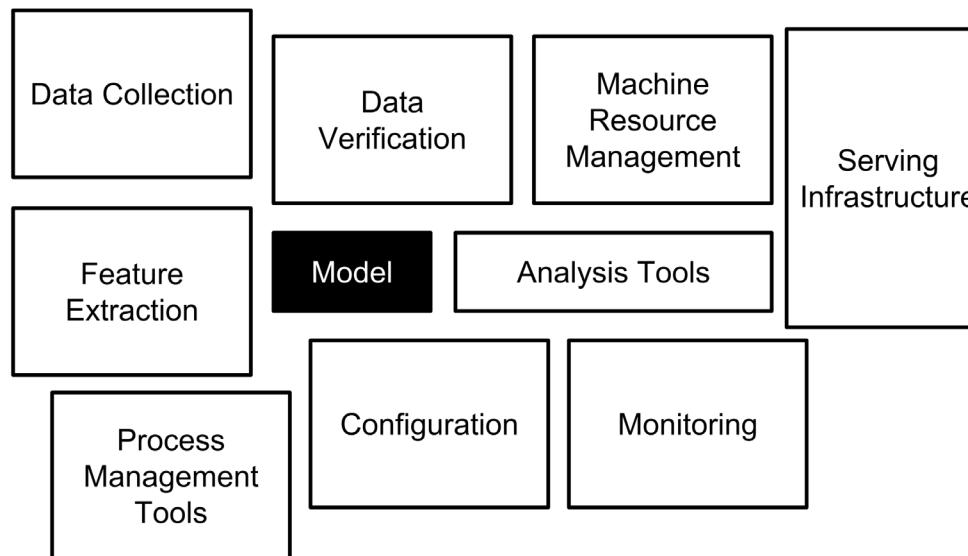
## *Part 1*

DIT826

Daniel Strüber

# Learning objectives

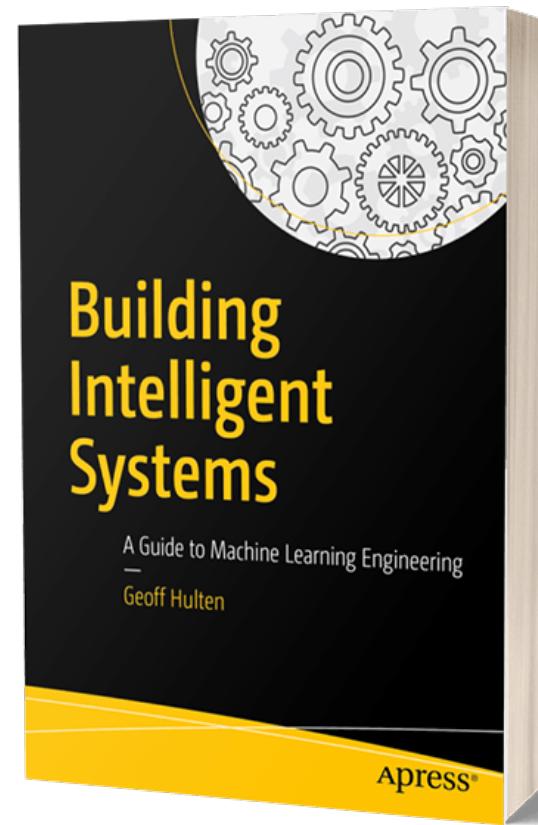
- **Understand** that building ML systems is more than training a model
- **Understand** practices and challenges of ML systems engineering



# Book

*Building Intelligent Systems:  
A guide to Machine Learning  
Engineering*, Geoff Hulten

Available as an e-Book at GU and Chalmers libraries

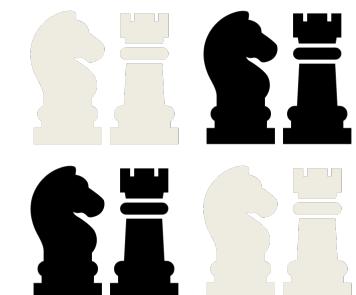


# Mandatory reading

- D. Sculley et al., Hidden Technical Debt in Machine Learning Systems, <https://papers.nips.cc/paper/5656-hidden-technical-debt-in-machine-learning-systems.pdf>
- S. Amershi et al., Software Engineering for Machine Learning: A Case Study, [https://www.microsoft.com/en-us/research/uploads/prod/2019/03/amershi-icse-2019\\_Software\\_Engineering\\_for\\_Machine\\_Learning.pdf](https://www.microsoft.com/en-us/research/uploads/prod/2019/03/amershi-icse-2019_Software_Engineering_for_Machine_Learning.pdf)
- M. Zinkevich, Rules of Machine Learning: Best Practices for ML Engineering,  
[http://martin.zinkevich.org/rules\\_of\\_ml/rules\\_of\\_ml.pdf](http://martin.zinkevich.org/rules_of_ml/rules_of_ml.pdf)

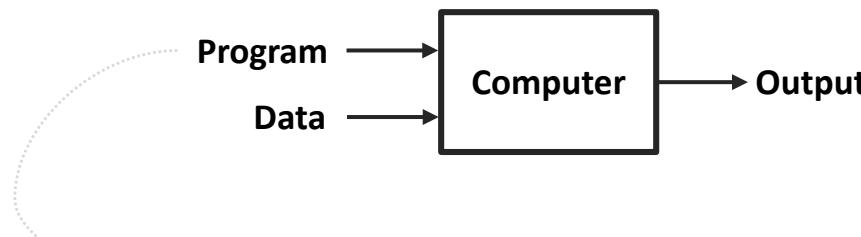
# Machine Learning

- Is a branch of artificial intelligence based on the idea that systems can learn from data, identify patterns, and make decisions with minimal human intervention.
- “A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its **performance at tasks** in **T**, as measured by **P**, improves with **experience E**” Tom Mitchell
- Example
  - Task T, What is the task?
  - Experience E, What is the Experience?
  - Performance P, What is the Performance?
    - *If P increases with E → the machine is learning!*



# Programming vs. ML

## Traditional Programming



## Machine Learning



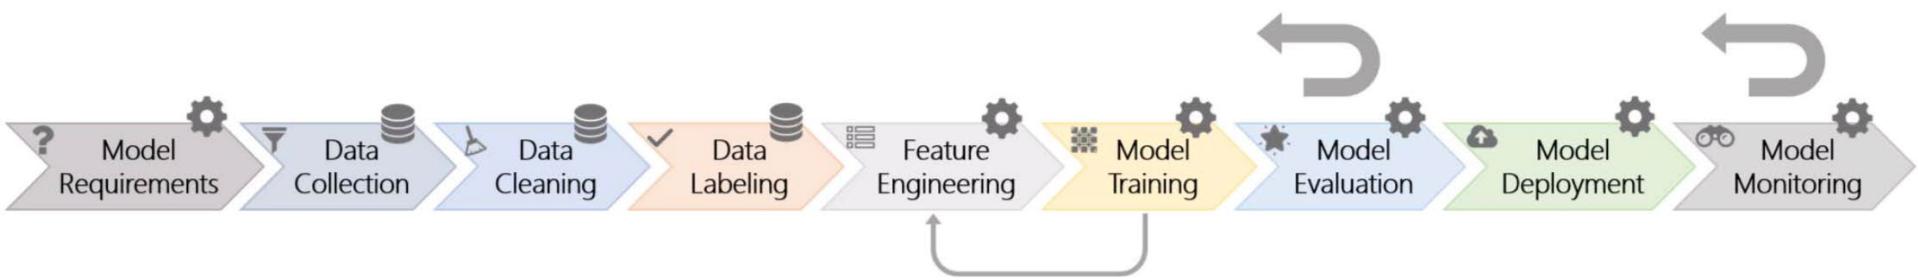
# Software Engineering

- A systematic approach for software development.
- An iterative process that includes different activities, such as:
  - planning, risk analysis, requirement engineering, software design, coding, versioning, testing, integration, deployment, maintenance, etc.

# Software Engineering For ML

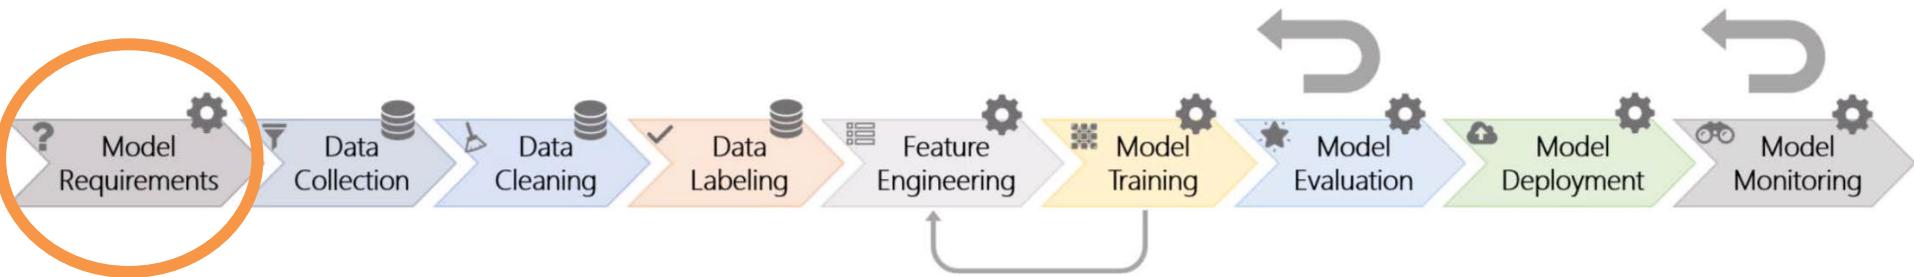
- A systematic way to create and integrate ML into software products and services.

# Workflow is needed



- Not a linear process

# The ML Workflow



- Model Requirements:
  - Problem and Goals?
  - Models?

# Problems and when to use ML

- Syndrome: “*We need ML because it’s trendy*”\*
- How **often** you think you need to **update** your system?
  - If ***n*** is small, then using ML is probably not right.
- Example:  
 $NewBalance = OldBalance - WithdrawalAmount$

\***Rule #1:** Don’t be afraid to launch a product without machine learning, in “M. Zinkevich, Rules of Machine Learning”

# Requirements: Problem

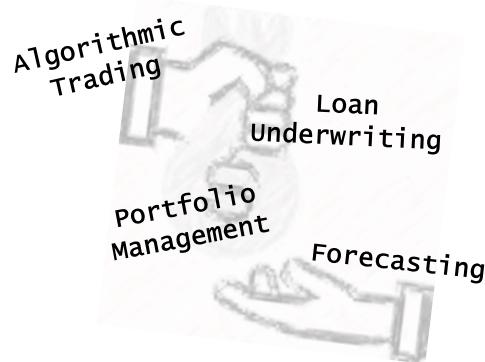
- **Big** problems
  - ~100 million songs
  - ~130 million books
  - ~1.5 billion websites
- **Time changing** problems
  - New technologies
  - Human faces – masks, face tattoos?
  - UX
- **Open ended** problems
  - ~6k tweets per second
  - ~60k new web pages per day
  - ~3 billion active Facebook users
- **Intrinsically hard** problems
  - Weather prediction
  - Complex Open-ended games
    - First chess program 1957
    - Deep Blue 1997 beat Kasparov
    - AlphaGo

# Success of ML and AI

Source: Geoff Hulten, *Building Intelligent Systems*



Web Search



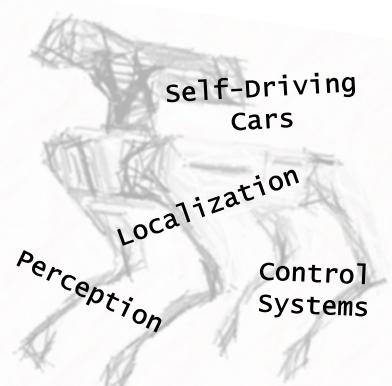
Finance



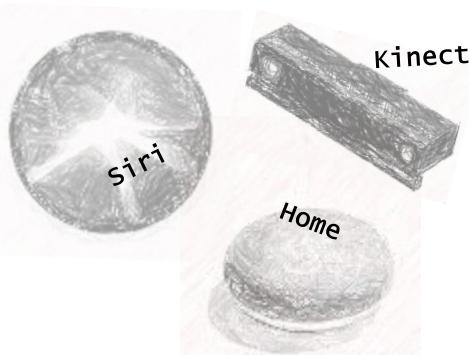
Marketing & E-commerce



Abuse / Security



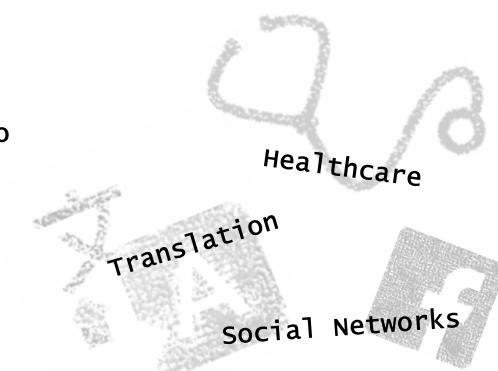
Robotics



Digital Assistants



Games



Many Others

# A Problem-Solving Technique: *Design Thinking*

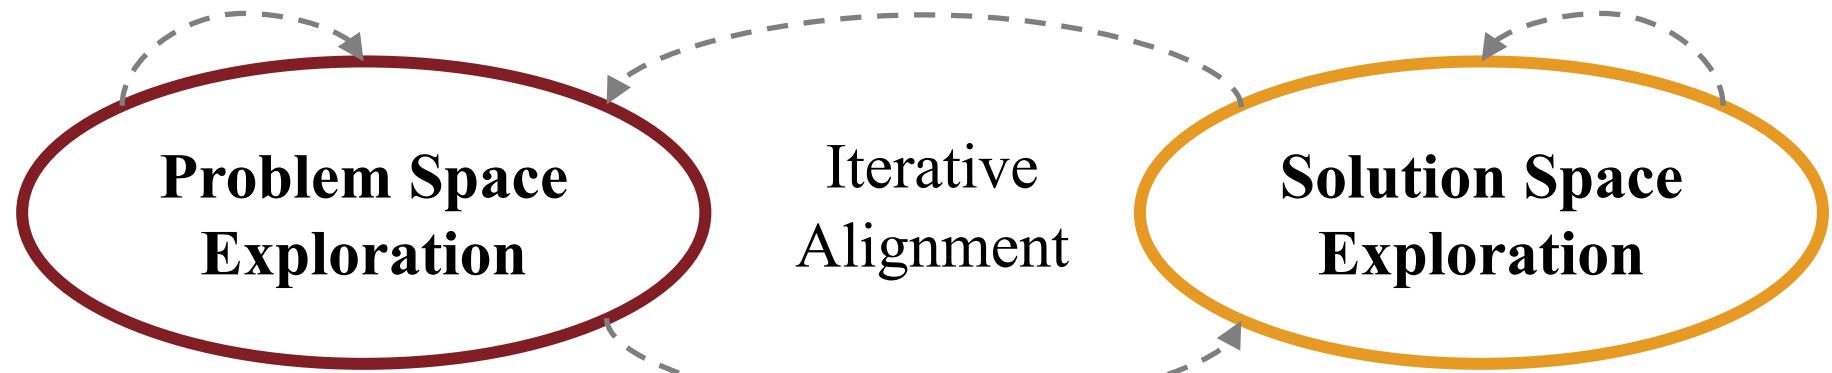
Framing and  
Analyzing

Ideating and  
Evaluating

**Problem Space  
Exploration**

Iterative  
Alignment

**Solution Space  
Exploration**



- Lindberg et al. *Design thinking: A fruitful concept for IT development?* In *Design thinking*. Springer, 3–18, 2011
- Nigel Cross. *Design thinking: Understanding how designers think and work*. Berg, 2011

# Defining the Goals

- A successful goal should:
  1. Communicate the **desired outcome**: *what?*
  2. Be **achievable**: *how?*
  3. Be **measurable**: *does it work?*
- Spam/Not Spam



# Model Type

Source: Geoff Hulten, *Building Intelligent Systems*

- What models are most appropriate for the given problem?  
(classification, clustering, etc.)

Decision Tree

Book Title	Number of Pages	Year Published	Genre	HasWord(Robot)	Author ID	Best Seller
Gone With The Wind	1037	1936	Historical Romance	0	1001	1
For Whom the Bell Tolls	480	1940	War Drama	0	1010	1
I, Robot	253	1980	Science Fiction	1	1020	1
One Hundred Goodbyes	100	2018	Science Fiction	0	1030	0

```
def F(BookTitle, NumberOfPages, YearPublished, Genre, HasWord(Robot), AuthorID):  
  
    if YearPublished > 1990:  
        if Genre == "Science Fiction":  
            return 1  
        else:  
            return 0  
    elif AuthorID == 1010:  
        return 1  
    else:  
        return 0
```

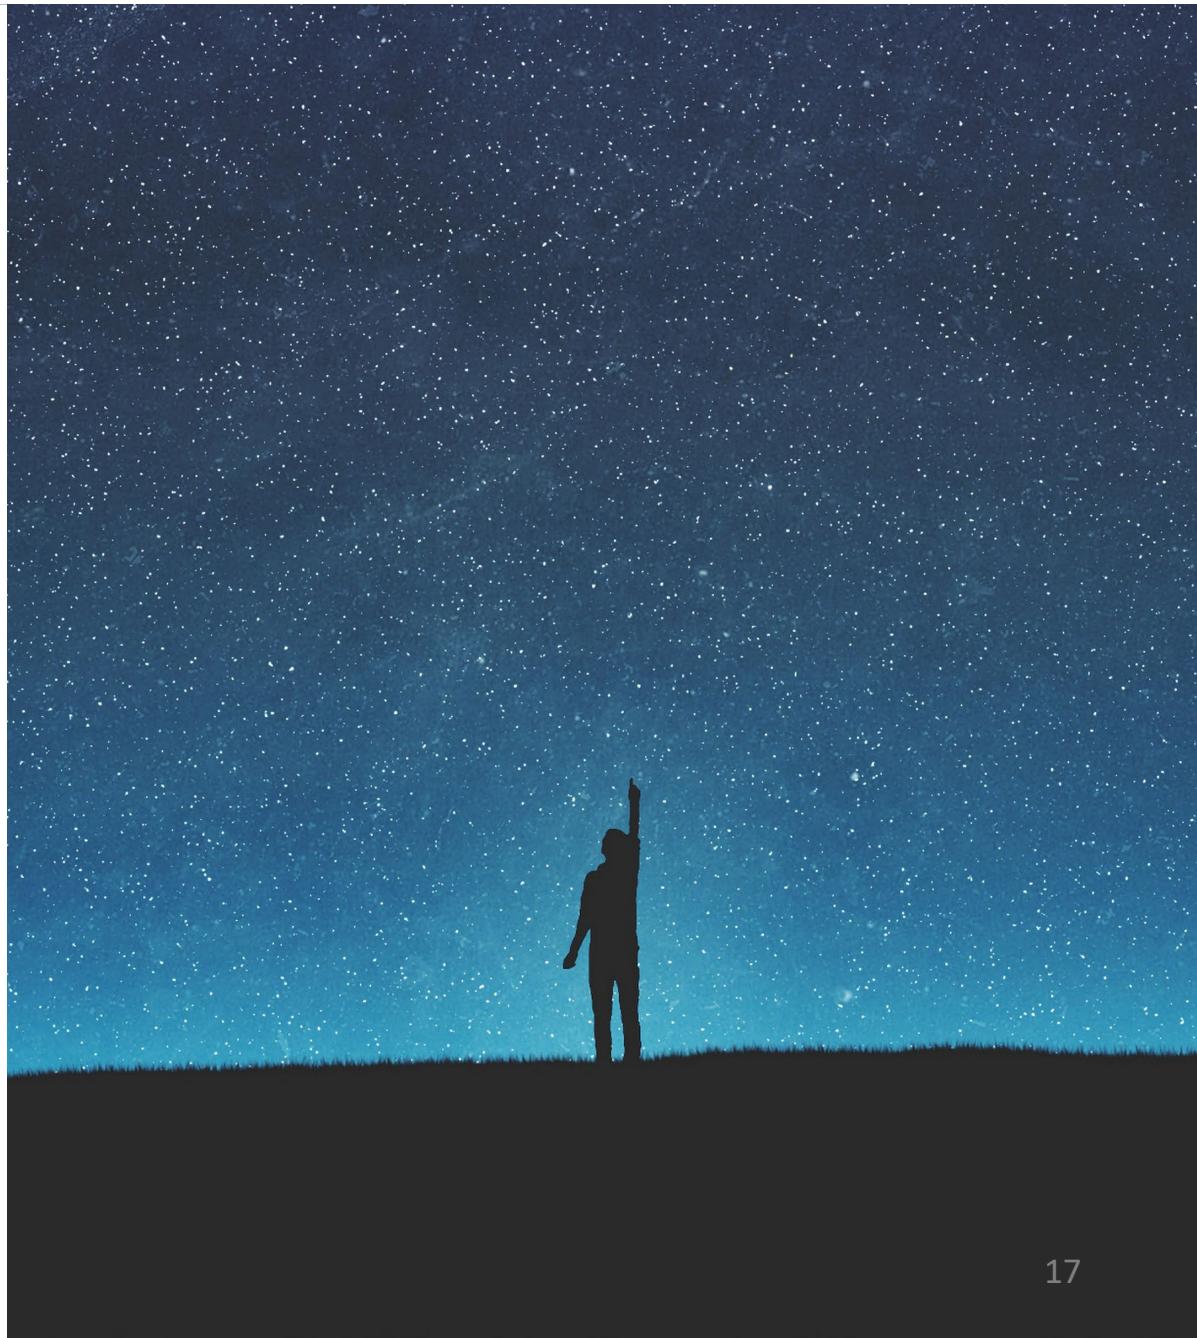
Linear Models  
Decision trees  
Ensembles of models  
Neural networks  
Support vector machines  
Etc.

F(X)
1
0
0
0

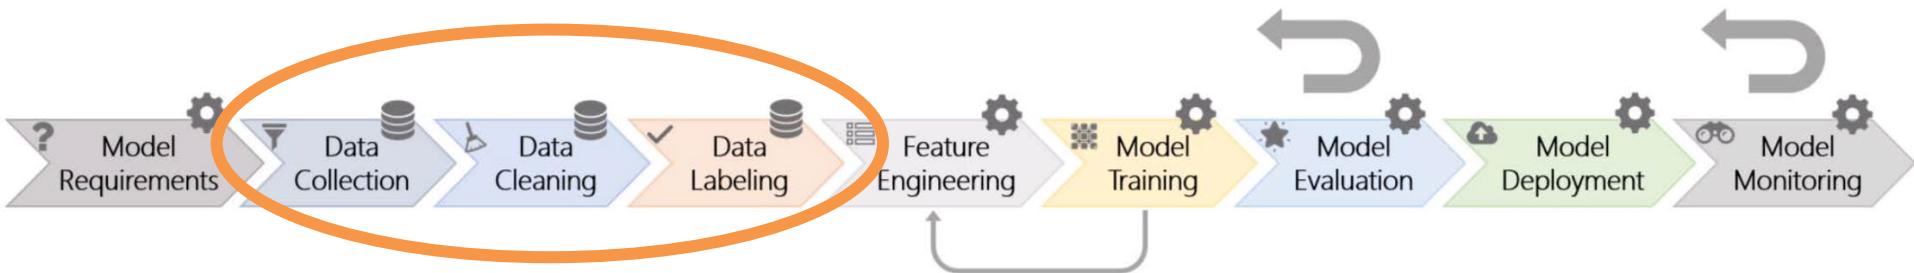
Linear Model

```
def F(BookTitle, NumberOfPages, YearPublished, Genre, HasWord(Robot), AuthorID):  
  
    sum = 0.5 * NumberOfPages + 0.75 * YearPublished + 0.1 * AuthorID  
  
    return 1 if sum > 2000 else 0
```

# Questions?



# The ML Workflow



- Working with data:
  - Collection
  - Cleaning
  - Labeling

# Data pipeline

## Definition

- Process that takes input data through a series of transformation stages, producing data as output.
- Both the input and output data can be fetched and stored in different locations, such as a database, a stream, a file, etc.

# Data collection

- Prepare the **scripts** to fetch the raw data
- **Store** the data

- CSV



- DB



- Cloud



# Data collection

- Sources
  - Existing data sources in the company
  - Existing open source data
  - Collect new data

The Kaggle logo, written in a large, blue, lowercase sans-serif font.

# Data cleaning

- The data are often **unstructured** and can be quite difficult to work with.
- Dealing with **noise** e.g., inaccurate and incomplete data.

# Data labeling

- Ground truth
- Sometimes readily available  
(sale price of house collected from a website)

# Pitfalls of working with data

## Broken confidence intervals

- 95% chance of being **within** an interval means that there is a 5% chance of being **outside** the interval.
- *Example:* Self-driving cars
  - 100 times → 5 mistakes
- TIP: Ask yourself few questions:
  - Is this right?
  - How sure are we?
  - Is there another interpretation?
  - How can we know which is correct?

# Pitfalls of working with data

## Noisy data

- Every large data set will have **noise**.
- Noisy data will inject **errors** into things created from these data.
- *Mitigation:*
  - Need for validation (There is lecture on this soon!)

# Pitfalls of working with data

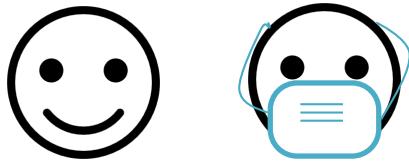
## Biased data

- Bias happens when data is collected in ways that are **systematically different** from the way the data is used.
- Bias can make data **less useful**.
- *Mitigation:*
  - Get more relevant telemetry or training data that contain context.

# Pitfalls of working with data

## Out-of-date data

- Things **do change** and collected data might not be representative anymore.
- *Example:* face recognition



- *Mitigation*
  - Train and deploy new models.

# Questions?

