

Engineering of ML Systems

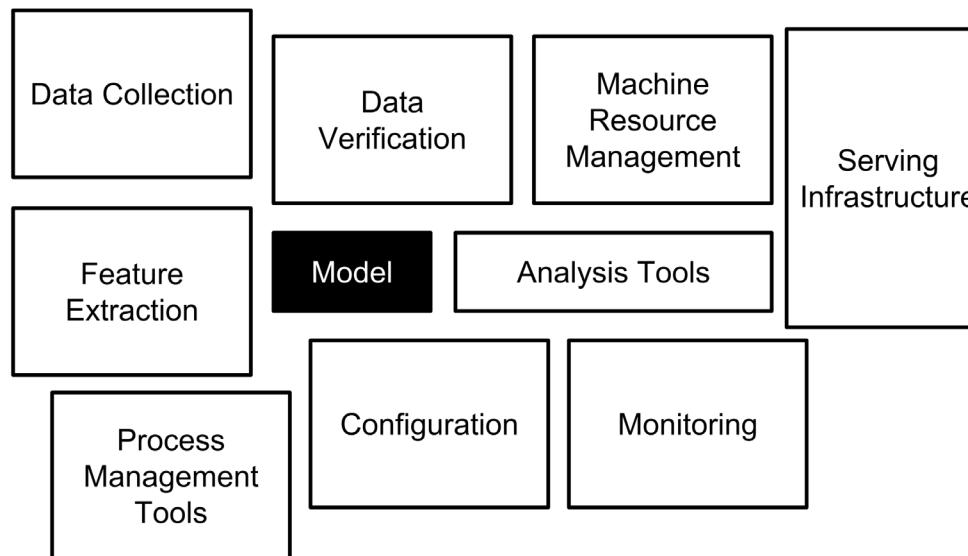
Part 2

DIT826


Daniel Strüber

Learning objectives

- **Understand** that building ML systems is more than training a model
- **Understand** practices and challenges of ML systems engineering







The ML Workflow



- Working with data:
 - More issues
 - Mitigations

TECH DEBT



SW engineering issues → data

Unstable dependencies

- SW design principle: “*minimize dependency to unstable modules*”
- Unstable data dependencies
 - Quality or quantity of data source deteriorates
 - Format or schema of data sources changes
- Especially problematic when models are retrained frequently

SW engineering issues → data

Unstable dependencies

- **Mitigation:** Create *versioned copy* (i.e., decoupling)
- **Mitigation:** Monitor for upstream instability in features
 - What alert would fire if one datacenter stops sending data?
 - What if an upstream signal provider did a major version upgrade?

CSV validator

Validation of CSV file

- <https://pypi.org/project/csvvalidator>

```
# import everything from the csvvalidator package
from csvvalidator import *

# Specify which fields (columns) your .csv needs to have
# You should include all fields you use in your dashboard
field_names = ('date',
               'units_sold',
               'store'
               )
# create a validator object using these fields
validator = CSVvalidator(field_names)
```

CSV validator

```
# write some checks to make sure specific fields
# are the way we expect them to be

validator.add_value_check(
    'date', # check for a date with the specified format
    datetime_string('%Y-%m-%d'),
    'EX1', # code for exception
    'invalid date' # message to report if error thrown
)
validator.add_value_check(
    'units_sold',
    int,
    'EX2',
    'number of units sold not an integer'
)
validator.add_value_check(
    'store',
    enumeration('store1', 'store2'),
    'EX4',
    'store name not recognized'
)
problems = validator.validate(data)
```

SW engineering issues → data

Unnecessary dependencies

- SW design principle: “*keep it simple*”
- Unnecessary data dependencies
 - Correlated features
 - ϵ -features (small improvement)
 - Forgotten features (from previous experiment)*

*Rule #22 – Clean up features you are no longer using, in “M. Zinkevich, Rules of Machine Learning”

SW engineering issues → data


Unnecessary dependencies

- Issue of **cost**
 - **Storage** cost
 - (esp. if data is massively large)
 - (esp. if data is versioned)
 - **Computational** cost
 - **Maintenance** cost
 - **Monitoring** cost

SW engineering issues → data

Unnecessary dependencies


- One Solution:
 - Feature selection via exhaustive '*leave-one-feature-out*' evaluations.




SW engineering issues → data

- SW design principle: “*avoid module dependency loops*”
- Hidden **feedback loops**
 - Output of model A → fed to model B
 - Output of model B → fed *back* to model A
- **Mitigation**
 - Reduce dependency
 - Better understand the problem that you want to solve
 - Testing

Questions?



The ML Workflow



- Working with data:
 - **Feature Engineering**

Feature Engineering

- The art of converting context into features that *work well* for your problem and with your model structure.



Feature Engineering

- Normalization and standardization: changing numerical features so they are comparable

$$\text{Standardized } f = \frac{f - \mu}{\sigma}$$

← Mean
← Std. Deviation

- Expose hidden Information
- Expand the context

Feature Engineering

Feature Selection


- *Which features to use?*
 - Use features that will help your model to **generalize**.
- *How many features to use?*
 - Approaches:
 - Frequency
 - Accuracy
 - ...

Feature Selection

Frequency

- Pick top N most **common** features in the training set.

Feature with a count > 1000



Feature	Count
to	1745
you	1526
I	1369
a	1337
the	1007
and	758
in	400
...	...


Feature Selection

Accuracy


- Pick **N** that improves the accuracy of the model.
 - Maximize the accuracy gain from using these features.
- Greedy search, **adding/removing** features **into/from** the model:
 - Add (**remove**) a candidate feature **into/from** the model.
 - Build and train the model and evaluate it.
 - Check the accuracy.
 - Add (**remove**) the feature.
 - Repeat until you get **n**.

Remove	Accuracy
<None>	88.2%
to	92.5%
FREE	85.8%
...	...

Questions?



The ML Workflow



- Working with data:
 - Model Training

A Challenge...

- Different **expertise** is necessary
 - Software development
 - Data management (large data)
 - Machine learning
- Different teams might work in **silos**
 - **Data engineers**: building data pipeline
 - **Data scientist**: selecting and tuning the algorithm
 - **Web developers**: building the client/server

...consequences

- Models might only **work in the lab**
 - Never leave the *proof-of-concept* state
- **Hard** to update and maintain.
- Many are advocating for **CD practices** applied to production ML systems*

* D. Sato, A. Wider, C. Windheuser, Continuous delivery for machine learning,
<https://martinfowler.com/articles/cd4ml.html>

Static vs. dynamic training

Definition

- **Static training**
 - Inference model trained **once** on a fixed data set
 - Put in production, does not change
- **Dynamic training**
 - Labeled data **continuously** coming into the system
 - New data incorporated into the model

Static vs. dynamic training

When to use

- **Static training**
 - Modeled reality is not expected to change
(e.g., cat/dog pictures recognition, playing chess)
- **Dynamic training**
 - Modeled reality has trends and changes
(e.g., buying trends, recommendations, self-driving cars, etc.)

Static vs. dynamic training

Attention points

- **Static training**
 - Model can degrade slowly.
 - More effort on data collection and labeling
- **Dynamic training**
 - Need monitoring (going off the rail?)
 - Need infrastructure for versioning and roll-back
 - Less effort on data collection, but tricky to get it right.

Reproducibility and auditability

- Keeping track of history is important
- To have **reproducible results**, need to know the **exact version** of
 1. Source code
(e.g. model training and pre-processing)
 2. Training data
(e.g. input signals and labels)
 3. Platform
(e.g. OS, GPU, and version of installed packages)

Traceability and versioning

- Data sets should be tagged with metadata
 1. Origin
 2. Freshness (when collected)
 3. Code used to extract it




Traceability and versioning

- Models should be tagged with provenance
 - Which data used for training/testing?
 - Which pipeline generated it?
- Versioned copies should be kept of each <code,data,model,environment> tuple
- At least for models that make it to production

Versioning challenges

- Versioning very large artifacts (e.g., data) is difficult and costly.

Questions?



What is different with ML?

Traditional
software project

A **bit** of this

A **lot** of this



Data



Model



Code

ML
software project

A **lot** of this

An this

A **bit** of this

Code is code, right?

Design anti-patterns in ML systems

- Pipeline jungles
 - no modularization of pipeline code, no refactoring
- *Mitigation*
 - Modularize and organize the code.
 - Testing.

Code is code, right?

Design anti-patterns in ML systems

- Excessive glue code
 - Massive amount of code written to get data into/out of a general-purpose package
- *Mitigation*
 - Create clean native solutions where feasible
 - Carefully consider the benefit and cost of adding (yet another) general-purpose package

Code is code, right?

Design anti-patterns in ML systems

- Dead experimental code paths
 - no clean up discipline
- *Mitigation*
 - Periodical inspection of experimental code.

Example

Blob of code. What does it do?

```
df = pd.DataFrame(...)

del df['column1']

df = df.dropna(subset=['column2', 'column3'])

df = df.rename({'column2': 'unicorns', 'column3': 'dragons'})

df['new_column'] = ['iterable', 'of', 'items']

df.reset_index(inplace=True, drop=True)
```


Example

Blob of code. Explained!


```
# create a pandas DataFrame somehow
df = pd.DataFrame(...)
# delete a column from the dataframe
del df['column1']
# drop rows that have empty values in column 2 and 3
df = df.dropna(subset=['column2', 'column3'])
# rename column2 and column3
df = df.rename({'column2': 'unicorns', 'column3': 'dragons'})
# add a new column
df['new_column'] = ['iterable', 'of', 'items']
# reset index to account for the missing row we removed above
df.reset_index(inplace=True, drop=True)
```

Good comments explain rationale

Questions?




The ML Workflow



- Working with data:
 - Evaluation is creation


Evaluation

Machine Learning



- Does **Process 1** do a good job at mapping ‘data’ to ‘output’?
- Is **Model 2** better than **Model 1**?
- Are the mistakes similar or different? Which is better?

The ML Workflow



- Where intelligence lives
- Intelligent experience
- Related to **software architecture** –
solutions are instances of architectural styles

Where Intelligence Lives

- Static intelligence in the product
- Client-side intelligence
- Server-side intelligence
- Back-end cached intelligence

Where Intelligence Lives

Static intelligence in the product

- Model is **packaged** with the application
- Pros
 - Cost of operation: Cheap
 - Latency in execution: Excellent
 - Offline operation: Yes
- Cons
 - Latency in updating intelligence: Poor
 - No data to improve the intelligence

Where Intelligence Lives

Client-side intelligence

- Executes **completely** on the client.
- Pros
 - Latency in execution: Excellent
 - Offline operation: Yes
- *It depends*
 - Latency in updating intelligence: Variable
 - Cost of operation: Based on update rate
- Cons
 - Exposes intelligence to the world

Where Intelligence Lives

Server-centric intelligence

- The intelligence runs in **real-time in a service**
 - Client sends features to a server; the server executes the intelligence on the features and returns the result.
- Pros
 - Latency in updating intelligence: Good
 - Easier monitoring
- It *depends*
 - Latency in execution: Variable
 - Cost of operation: Variable
- Cons
 - Offline operation: No

Where Intelligence Lives

Back-end cached intelligence

- Involves running the intelligence off-line, caching the results, and delivering these when needed.
- Pros and Cons
 - Latency in execution: Variable
 - Latency in updating intelligence: Variable
 - Cost of operation: Variable
 - Offline operation: Partially


Intelligence Experience

- A **connection** between the intelligence and users.
- *Example*: an intelligence determines that the user is in a room where it is a little dark for humans to see comfortably.
 - How should the experience respond?

Intelligence Experience

Techniques

- $P(\text{LikeSong} \mid \text{User, History})$
- Modes of Intelligent Interaction



- Which intelligence experience to use? When to Update it?

Intelligence Experience

Goals

- Achieve objectives & increase profit and sales
- Engage users and achieve their objectives
- Get data to improve
- *Example:* Home Light → $P(\text{on} \mid \text{sensors, motion, etc.})$
 - Objective?

 $P \leq 0.50$ $P > 0.50$ $P > 0.90$; turn on light $P > 0.75$; then prompt user $P < 0.25$; turn light off

(turn light off after time t)



Full automation




Save Energy




User Safety

Questions?



The ML Workflow



- Model Monitoring

Monitoring and Telemetry

- **Monitoring** is foundational for producing intelligence that:
 - functions correctly, and
 - improves over time.
- It includes knowing:
 - Context
 - Answers
 - Experiences
 - User behavior
- **Telemetry**: collect data at remote points and transmit them to a central point for monitoring



Monitoring inference output


- Model output
 - Predicted value is within bounds
 - Warning sign of model going off the rail
- Add sanity checks on output
- Examples
 - Predicted house not 10x bigger in same area
 - Values are within specified distributions

Monitoring: Rolling back

- Model “roll back” procedure: Being able to quickly revert to a **previous known-good model**
- **Test** how quickly and safely a model can be rolled back.

Intelligence Orchestration

- Achieve **objectives** of the ML system
- Controls the **monitoring**
- Balance **experience**
- Deals with **mistakes**




Orchestrating intelligent systems


Why needed?

- **Objective changes:**
 - Better understand a problem
 - Solved previous objective
- **Intelligence changes:**
 - More data, better models.
 - More accurate or less accurate model.
- **User changes:**
 - New users or users leave
- **Cost changes:**
 - Telemetry costs
 - Mistakes costs

Questions?



Summary: lifecycle of data-intensive AI systems




Prepare data

- Ingestion, Labeling, Normalization, Transformation, Validation

Source: D. Sato, A. Wider, C. Windheuser, Continuous Delivery for Machine Learning
<https://martinfowler.com/articles/cd4ml.html>


Lifecycle of data-intensive AI systems



Experiment and build model

- Select features, Select ML algorithm, Tune parameters, perform evaluation and testing

Lifecycle of data-intensive AI systems



Serving the model in production

- Deploy, Predict, Monitor, Update model

Questions?

